

Please use the following QR code to check in and record your attendance.

00:01:59

CS 1037

Fundamentals of Computer  
Science II

# Multi-way Search Tree

---

Ahmed Ibrahim




# Recall: Binary Search Trees (BST)

- What is a Binary Search tree?


A binary search tree is a **binary tree** in which every node contains only smaller values in its left subtree and only larger values in its right subtree.

- Every BST is a BT, but every BT must not be a BST.
- There must be no duplicate nodes (in general).



# BTs Drawbacks



- In a BST, each node can have **only two children**, so the tree's height grows quickly as more nodes are added. For large datasets, this height increase makes searches slower, as it takes more steps to reach a **leaf** from the **root**.
- **Balance** is essential for good performance in a BST. If nodes are inserted in a sorted order (like ascending or descending), the tree can easily become unbalanced. Self-balancing trees (like AVL) solve this issue but add extra complexity and require rebalancing.



# The concept of Multi-way Search Trees

- A multi-way tree is a tree that can have more than **two children**.
- A multi-way tree of order  $m$  (or an  $m$ -way tree) is one in which a tree can have  $m$  children.

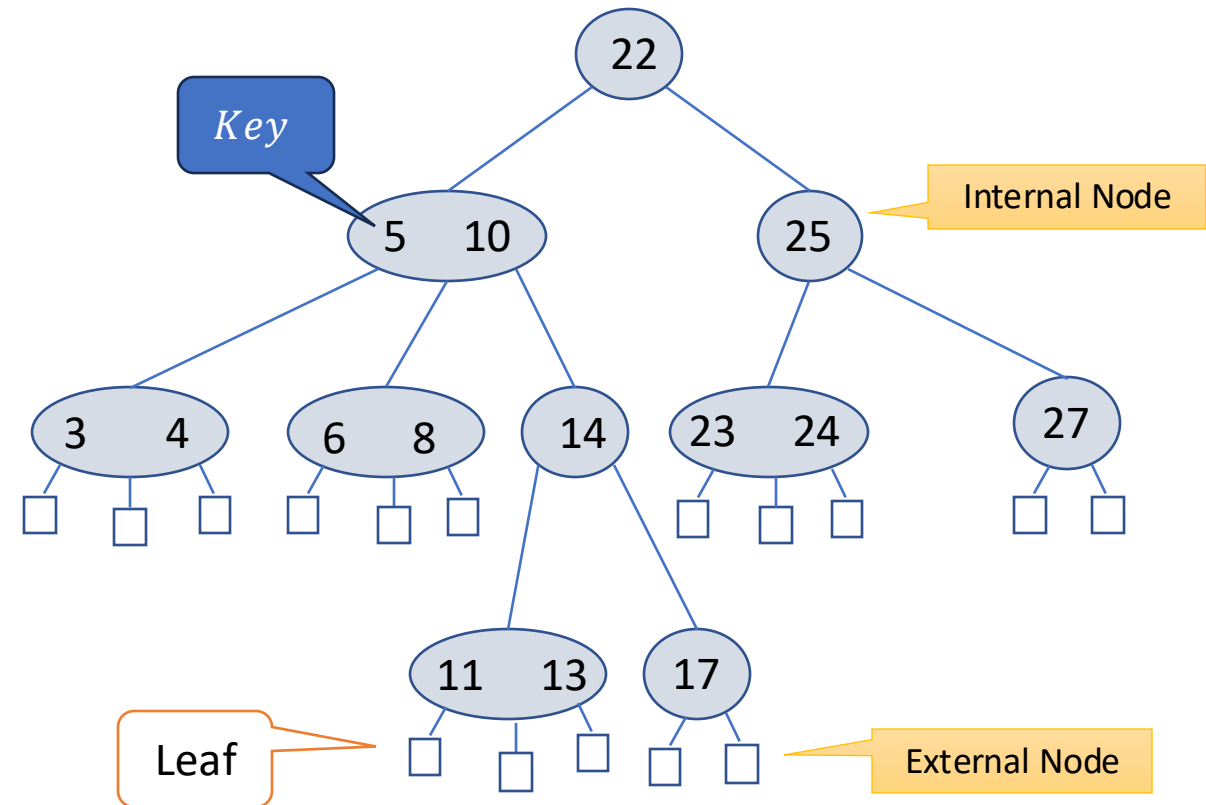

- Legend:



Example of 3-way tree

# Multi-way Search Tree

- As with the other studied trees, the nodes in an  $m$ -way tree consist of  $m-1$  **entries** and pointers to children.
- The **entries** are in the form of pairs  $(k,D)$  where  $k$  is the **key** and  $D$  is the value (**data**) associated with the **key**.
- The external nodes of a  $m$ -way search tree do not store any **entries** and are “*dummy*” nodes.



# The Structure of Multi-Way Search Trees

The following structure defines a simple  $m$ -way node structure, where  $m$  ( $>1$ ) is a predefined constant.



```
typedef struct node {  
    int count; // number of key values  
    int key[m-1]; // key arrays  
    struct node *child[m]; // sub-tree pointer array  
} TNODE;
```

# Properties of $m$ -Way Search Trees

- A  $m$ -way search tree is an **ordered tree** such that
  - Each **internal node** has at least two and at most  $m$  children and stores  $m-1$  data items
  - **External nodes** have zero data items
  - Number of children = 1 + number of data items in a node
    - A node with three children is called a 3-node.





# An Internal Node



Numerical Example

- Each internal node has  $m \geq 2$  children and stores  $m-1$  entries.


# Multi-Way Tree Traversal



- In a multi-way search tree, in-order traversal involves visiting all keys in sorted order by **recursively** traversing child nodes and printing keys. Unlike BST, multi-way search trees may have more than two children per node, so the traversal is slightly different.
- An **in-order** traversal of a multi-way search tree visits the keys in **increasing order**.

# Multi-Way Tree Traversal


- Traversal in a BST visits each node about its two children (e.g., left subtree → node → right subtree in in-order).
- Traversal in a multi-way search tree involves visiting each key in the node and then recursively traversing each child subtree in the appropriate order. For example:
- In-order traversal in a multi-way tree involves:
  1. Recursively traversing the first child.
  2. Visiting the first key.
  3. Recursively traversing the second child.
  4. Visiting the second key.
  5. And so on for all keys and children.



# Multi-Way Tree Traversal Algorithm

- The following program is an example of in-order traversal with printing key values.


```
/* in-order traversal of m-way tree*/  
void print_inorder(TNODE *root) {  
    if (root != NULL) {  
        // Traverse the first child subtree  
        print_inorder(child[0]);  
        int i;  
        for (i=0; i < root->count; i++) // Traverse through each key in the node  
        {  
            printf("%d ", root->key[i]); // Print the current key  
            print_inorder(child[i+1]); // Traverse the next child subtree  
        }  
    }  
}
```




# Example

---

- Consider a 3-way search tree with the following structure:
- **In-order traversal:**
- Traverse the leftmost child [5] → Visit key 10 → Traverse child [15] → Visit key 20 → Traverse child [25, 30].
- Result=> [5, 10, 15, 20, 25, 30].




# Multi-Way Search



- Assuming ***m*** is a constant independent of a number of nodes, examining each node takes constant time ( $O(1)$ ); thus, the time to search is **proportional** to the tree **height (*h*)**.
- Within each node, searching among the keys takes  $O(1)$  time if the number of keys (*m*) is constant.
- Imagine that we are searching for  $k = 30$

# Another Example: Multi-Way Searching

---



- Search for key 7
  - Search terminates at a leaf child, which implies there is no entry with key 7

# Multi-Way Searching Pseudocode

**Algorithm** get(r,k)

**In:** Root r of a multiway search tree, key k

**Out:** data for key k or null if k not in tree

**if** r is a leaf **then return** null

**else {**


    Use binary search to find the index i such that either

- $r.keys[i] = k$ , or
- $r.keys[i] < k < r.keys[i+1]$

**if**  $k = r.keys[i]$  **then return**  $r.data[i]$

**else return** get(r.child[i],k)


**}**







# (2,4)-Tree

- (2,4) tree is a special ***m*-way search tree** with the following properties:
  - **Node-size property** – every internal node has at most **four children**. Every node can store between 1 and 3 keys.
  - **Depth property** – all **external nodes** have the **same depth** (the tree is balanced).
- **Recall** that in a multi-way search tree, the minimum number of children for a node is 2
  - Thus, a node can have 2, 3, or 4 children; thus a **(2,4) tree** is also called a **2-3-4 tree**.
- Used in databases and filesystems due to their efficient **search, insertion, and deletion**.



# Insertion in (2,4) Trees

- How do we find the correct (preserving order) node  $v$  to insert?
- Case 1 – key  $k$  is already in the tree.**
  - Perform a search  $\Rightarrow O(\log n)$
  - when reached node  $v$  storing  $k$ , continue the search in the subtree to the of  $v$
  - stop when reaching the node with only leaf children
- Example: insert 15



# Insertion in (2,4) Trees

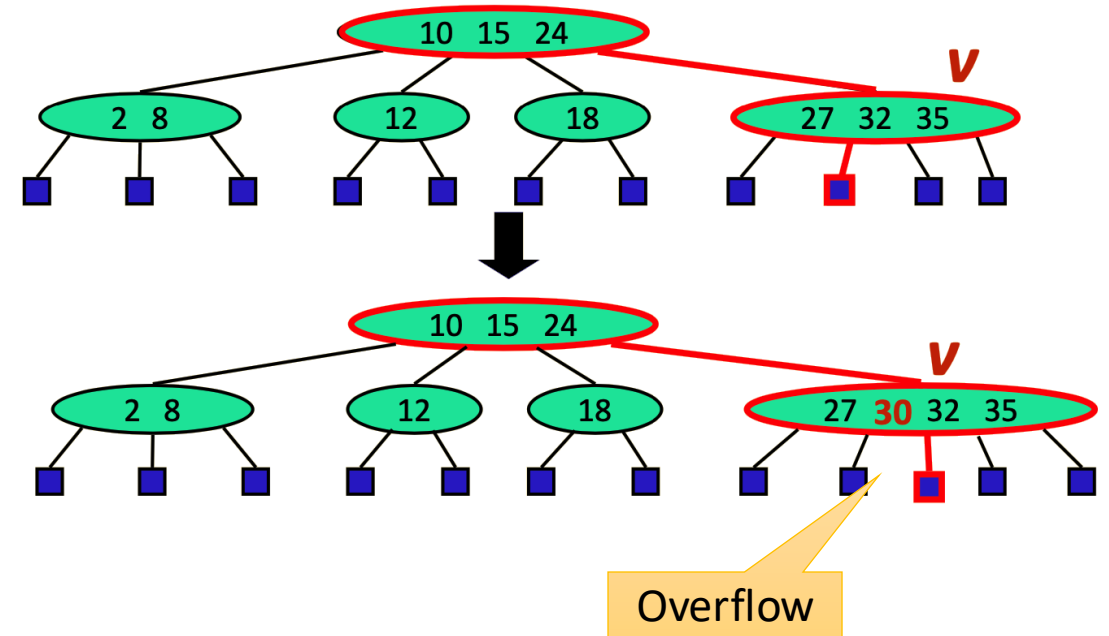
- How do we find the correct (preserving order) node  $v$  to insert?

- **Case 2 – Key  $k$  is not in the tree.**

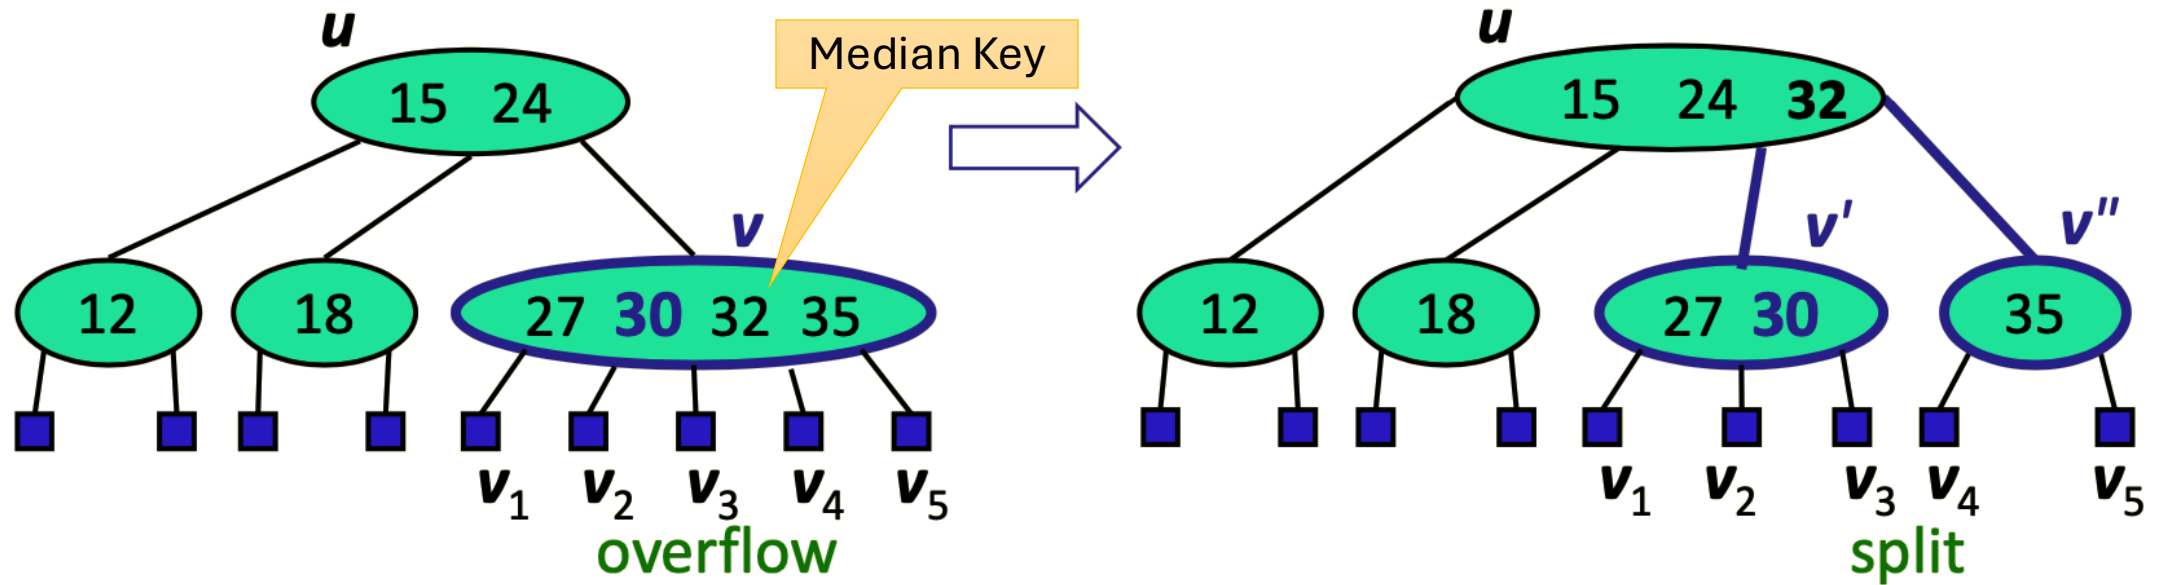
1. Perform a search  $\Rightarrow O(\log n)$
2. If key  $k$  is not in the tree
3. Then  $v$  is the leaf's parent reached by when searching for  $k$ .

- Example: insert 30
- However, there is a **violation** in the node size property

- **Overflow** occurs when a 4-node becomes a 5-node, illegal in (2,4)-tree






# Insertion: Overflow and Split



- The median key is the *third* key in the sorted list. It splits the  $v$  node into two smaller nodes.
- *Overflow* may propagate to the parent node  $u$ .

# Create (2,4) Trees




# Insertion in (2,4) Tree Algorithm

---

1. **Start at the root of the tree (T).**
2. **If the root is full (contains 3 keys):**
  - a. Split the root into two nodes.
  - b. Promote the *middle key* to create a new root.
  - c. Update the tree so the new root has two children.
3. **Find the correct leaf node where  $k$  belongs:**
  - a. Traverse down the tree starting from the root.
  - b. At each node, determine which child to move to based on the ranges defined by the keys.
  - c. Repeat until a leaf node is reached.
4. **Insert  $k$  into the leaf node:**
  - a. Add  $k$  to the leaf node's appropriate position (keys remain sorted).
5. **Check for *overflow*:**
  - a. If the node now contains more than 3 keys:
    - i. Split the node into two nodes.
    - ii. Promote the *middle key* to the parent node.
    - iii. Update the parent to reflect the new structure.
6. **Repeat the overflow process:**
  - a. If the parent node *overflows* after promotion, *split* the parent and promote its *middle key* to the next level-up.
  - b. Continue this process recursively up the tree until no overflow occurs or a new root is created.


# Deletion in (2,4) Trees

- *Case 1* – If an entry is an internal node with no leaf children, replace the entry with its **in-order successor** and delete the latter entry and one leaf.
- The **in-order successor** of a key in a (2,4)-tree is the smallest key larger than the given key, based on an in-order tree traversal.




# Underflow and Transfer

- *Case 2* – Assume that the entry to be deleted is at node  $v$  with leaf children.
- Deletion from  $v$  can cause an *underflow* (if  $v$  becomes a 1-node).
- To deal with *underflow* at node  $v$  with parent  $u$ , consider two cases:
  - **Transfer** operation, if an adjacent sibling  $w$  has at least **two entries** (3 leaf nodes)
  - **Fusion** operation, if all adjacent siblings of  $v$  are 2-node






# Transfer Operation



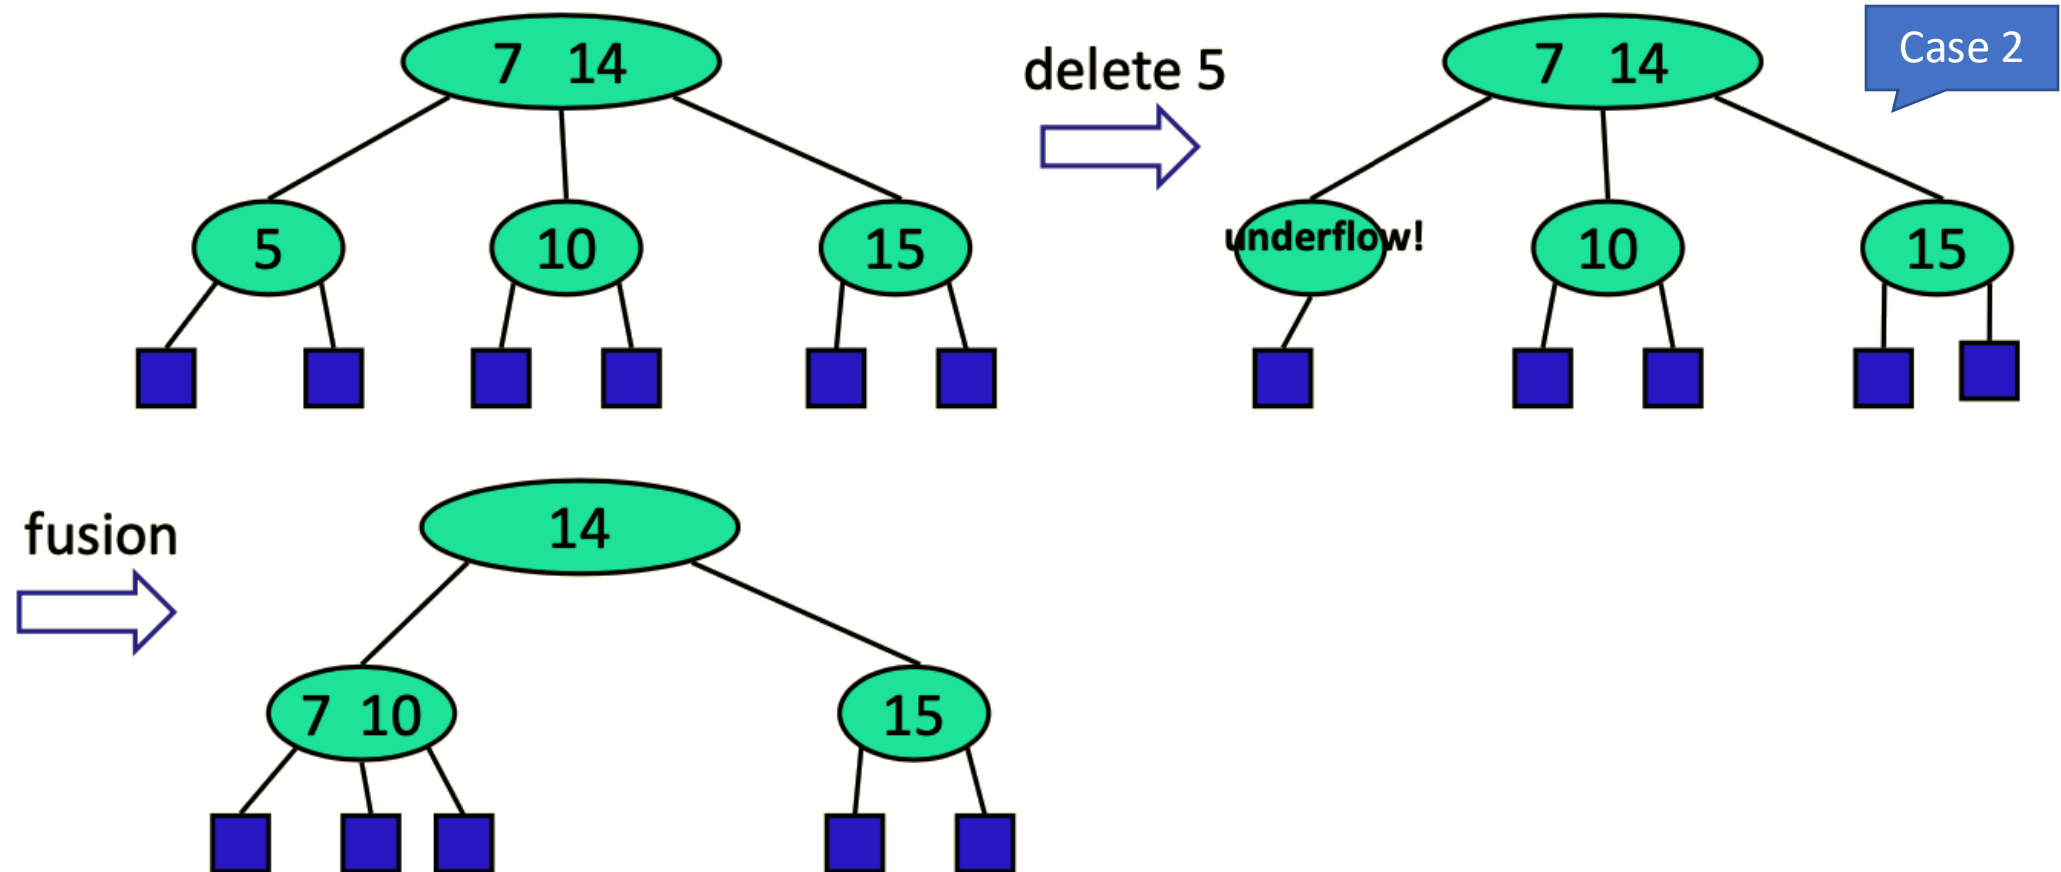
- An adjacent sibling  $w$  of  $v$  is a 3-node or a 4-node
  - Move an entry from  $u$  (the entry “between”  $w$  and  $v$ ) to  $v$
  - Move an entry from  $w$  (the entry with the key closest to the deleted key in  $u$ ) to replace the missing entry of  $u$

# Fusion Operation



- All adjacent siblings of  $v$  are 2-node
- Here, we are going to do a **fusion**. After a fusion, *underflow* may propagate to the parent  $u$ .

# Another Example



# Deletion in (2,4) Tree Algorithm

---

1. **Start at the root of the tree (T).**

2. **Search for key  $k$  in the tree:**

- a. Traverse down the tree, checking each node for  $k$ .
- b. If  $k$  is found in an **internal node**:
  - i. Replace  $k$  with its **in-order predecessor** (largest key in the left subtree).
  - ii. Recursively delete the predecessor key from the corresponding subtree.
- c. If  $k$  is found in a **leaf node**, proceed to step 3.

3. **Delete key  $k$  from the leaf node:**

- a. Remove  $k$  from the node.
- b. If the node still has at least 1 key, stop.
- c. If the node becomes *underfull* (0 keys), fix the deficiency described in step 4.

# Deletion in (2,4) Tree Algorithm

---

## 4. Fix *underfull* nodes (fewer than 1 key):

### a. Borrow from a sibling (Transfer):

- i. Check if a sibling node (adjacent child of the same parent) has more than 1 key.
- ii. Borrow a key from the sibling and move the corresponding parent's key into the deficient node.

### b. Merge with a sibling (Fusion):

- i. If no sibling has extra keys, merge the deficient node with a sibling.
- ii. Move the parent's key that separates the two siblings into the merged node.

## 5. Handle root *underflow*:

### a. If the root becomes *underfull* (0 keys) and has children:

- i. Promote the only child to become the new *root*.


### b. If the *root* is empty and has no children, the tree becomes empty.



# Hash Table

# Key-Value Entries

- You have at most  $N$  entries  $(k, v)$
- Suppose keys  $k$  are unique integers between 0 to  $N - 1$
- Create initially empty array  $A$  of size  $N$
- Store  $(k, v)$  in  $A[k]$
- Example



- Main operations (insert, find, remove) are  $O(1)$
- We need  $O(N)$  space

# Imagine that!

- What if we have 100 entries with integer keys, 0 to 1,000,000,000?
  - Do we still have  $O(1)$  insert(), delete(), find()?
  - We do not want 1,000,000,000 memory cells to store only 100 entries.
- What should we do?



# The Concept of Hash Table

- Array of Fixed Size (**TableSize**)
- Each key is mapped into some number between 0 and (**TableSize - 1**). Mapping is done by something called the **hash function**
- The **hash function** ensures that two distinct keys are assigned to different cells.
- Given the finite number of cells and an almost limitless supply of keys, a hash function is necessary to evenly distribute the keys among the cells!

|   |            |
|---|------------|
| 0 |            |
| 1 |            |
| 2 |            |
| 3 | john 25000 |
| 4 | phil 31250 |
| 5 |            |
| 6 | dave 27500 |
| 7 | mary 28200 |
| 8 |            |
| 9 |            |


# A Design Challenge

---


- Imagine a company has 5,000 employees and wants to store information about each employee using a hash table data structure.
- The company wants to use the employee's Social Security Number (SSN) of 9 digits as the key to looking up the corresponding employee information.
- The employee information includes the employee's **name**, **address**, **phone number**, and **salary**.
- The company needs to quickly look up an employee's information based on their **SSN**, as well as **add** and **remove** employees from the hash table.

# Possible Solution

- Hash table for storing entries (SSN, info)
- The hash function  $h(x) \Rightarrow$  last four digits of SSN



- Thus, an array of size  $N \Rightarrow 10,000$ 
  - The SSN is always of a fixed length.



# Hash Tables and Hash Functions


---

- A hash table is a data structure that allows for efficient storage and retrieval of **key-value pairs** using a **hash function**.




- **Hash Tables**
  - **provide** fast data retrieval and insertion, typically in constant time,  $O(1)$ , under ideal conditions.
  - **use** a hash function to transform input data (keys) into a fixed-size numerical value, determining where the data is stored in the table.

What if *keys*  
are NOT  
*integers*?



| Key     | Value |
|---------|-------|
| "Paul"  | 29    |
| "Jane"  | 35    |
| "Chloe" | 88    |
| "Alex"  | 18    |
|         |       |



# Hashing Non-Integer Keys


- Array **A** of size **N** = 8
- Design function  $h(k)$  that maps key **k** into integer range  $0, 1, \dots, N - 1$
- Entry with key **k** is stored at index  $h(k)$  in the array **A**

How far 'p' from 'a' => 15       $h(k) \Rightarrow 15 \bmod 8 = 7$

Distance = ASCII code of **k** - ASCII code of 'a'.

Hashing value

- ASCII code of 'p' = 112.
- ASCII code of 'a' = 97.
- Distance from 'p' to 'a':  $112 - 97 = 15$ .




| Key     | Value |
|---------|-------|
| "Paul"  | 29    |
| "Jane"  | 35    |
| "Chloe" | 88    |
| "Alex"  | 18    |

# Collision

- Collisions occur when different elements are mapped to the same cell.
- Collision resolution strategies
  - **Separate Chaining** – Store colliding keys in a **linked list** at the same hash table index
  - **Open Addressing** – Store colliding keys elsewhere on the table

# Collision Resolution by Chaining

- What if you still have  $N$  keys, which may not be unique? (1, A) (1, R) (1, C) (3, D)



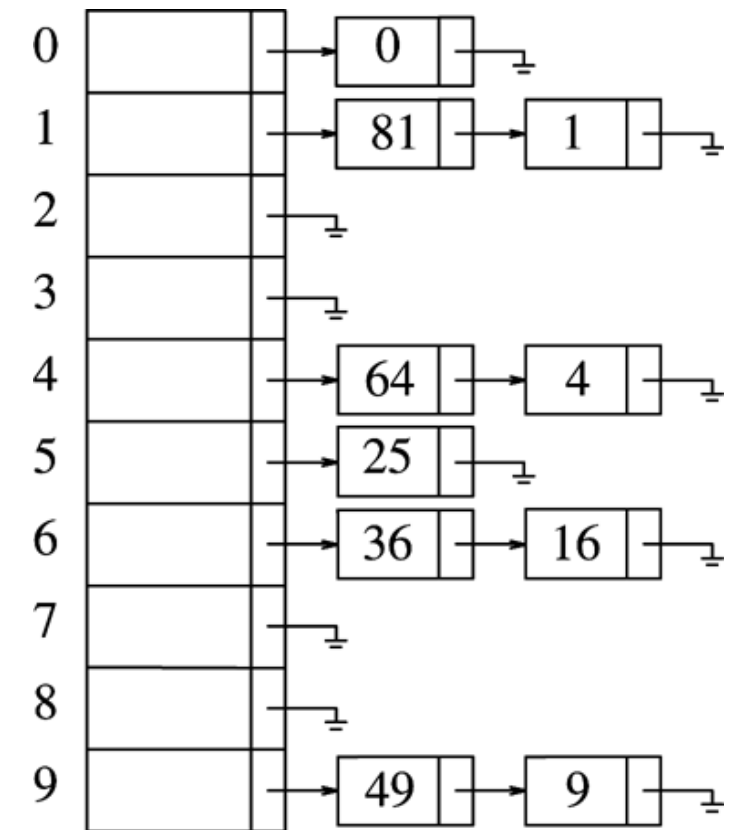
- A bucket array can be implemented as a **linked list**.
- Assume have at most a constant number of repeated keys methods `find()`, `remove()`, `insert()` are still  $O(1)$



# Example

- Hash table  $T$  is a vector of lists
  - Only singly linked lists are needed if memory is tight
- Key  $k$  is stored in the list at  $T[h(k)]$
- E.g. TableSize = 10
  - $h(k) = k \bmod 10$

Insertion sequence = 0, 1, 4, 9, 16, 25, 36, 49, 64, 81



# Load Factor $\lambda$

---

- The **load factor ( $\lambda$ )** is a measure that describes how **full a hash table** is.
- It is defined as:  $\lambda = N / M$  Where:
  - N: The total number of elements stored in the hash table.
  - M: The total number of slots in the hash table.
- The average length of a chain is equal to the **load factor**
  - A smaller load factor indicates fewer collisions and better performance.
  - A larger load factor increases the likelihood of collisions, leading to longer chains.
- Ideally, we want  $\lambda \leq 1$  (not a function of N)
- To maintain  $\lambda \leq 1$ , the hash table should resize (**rehash**) when it becomes **too full**.
- Keep the TableSize **prime** to ensure a good distribution

# Example

- Imagine a hash table with  $M=10$  (array slots) and  $N=7$  elements.
- The load factor is:  $\lambda = N / M = 9 / 10 = 0.9$
- We insert the following elements into the hash table= >  
Keys: 0, 81, 64, 25, 36, 49, 1, 4, 16
- With a hash function:  $h(k) = k \bmod M$  (where  $M=10$ ).
- **Collisions** may occur in buckets 0 and 6. We could consider resizing the table (rehashing) to reduce  $\lambda$  and minimize collision.

| Index | Keys Stored (Chaining) |
|-------|------------------------|
| 0     | 0, 81                  |
| 1     | 1                      |
| 2     |                        |
| 3     |                        |
| 4     | 4, 64                  |
| 5     | 25                     |
| 6     | 36, 16                 |
| 7     |                        |
| 8     |                        |
| 9     | 49                     |



Thank  
you